Culinary Computation

1st Nelson
Computer Science
Portland State University
Portland, United States
ajn6@pdx.edu

2nd Vassilevski
Mathematics
Portland State University
Portland, United States
panayot@pdx.edu

Abstract—Data utilization continues to become increasingly important as the rate of collection and transmission rapidly expand. Much work has been done investigating algorithmic creativity for music composition, visual artwork, and prose generation. This paper explores the potential of algorithmic generation and modification of recipes by describing the process of data collection, initial analysis, and processing technique.

Index Terms—graph, algebraic multigrid, data science, computational creativity

I. Introduction

Food is a central component of the human condition. Its expression acknowledges no barriers and is a defining component of all cultural experience. The US restaurant industry currently employs over 11 million individuals [1]. Understanding flavor theory and developing recipes isn't difficult for a trained professional but there is potential for improvement through utilization of data analysis. This is especially true for large corporate chains with extensive historical data.

II. Data Collection

Recipes were scraped from allrecipes and New York Times Cooking (NYTC). The allrecipes database is user created with around 60000 recipes. The NYTC database has around 18000 recipes and is curated and maintained by journalists. Another aggregated database curated by Majumder et al. [2] was also used.

III. Data Wrangling

The scraped recipes are in the form of a list of strings for ingredients and procedure. The quantity, ingredient, and adjectives were parsed from each ingredient in the ingredient list using a condition random field [3]. This process isn't perfect but yields high enough quality results to create the necessary data structures. Each unique ingredient in the database is given an ID and the recipes are converted into lists of IDs.

Using this new database an ingredient co-occurrence graph is constructed and represented as an adjacency matrix. This is a symmetric positive definite matrix with

a diagonal of 0s and each entry or weight is the number of times any two ingredients appeared in the same recipe.

For example, consider the following chicken recipes:

Recipe 1	Recipe 2	Recipe 3
chicken	chicken	chicken
lemon	salt	butter
salt	pepper	salt
pepper	garlic	pepper
olive oil	honey	garlic
oregano	water	paprika
parsley	rice vinegar	
	soy sauce	

All unique ingredients with IDs

0 — chicken	1 — lemon	2 — salt
3 — pepper	4 — olive oil	5 — oregano
6 — parsley	7 — garlic	8 — honey
9 — water	10 — rice vinegar	11 — soy sauce
12 — butter	13 — paprika	

Resulting co-occurrence matrix

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	-	1	3	3	1	1	1	2	1	1	1	1	1	1
1	1	-	1	1	1	1	1	-	-	-	-	-	-	-
2	3	1	-	3	1	1	1	2	1	1	1	1	1	1
3	3	1	3	-	1	1	1	2	1	1	1	1	1	1
4	1	1	1	1	-	1	1	-	-	-	-	-	-	-
5	1	1	1	1	1	-	1	-	-	-	-	-	-	-
6	1	1	1	1	1	1	-	-	-	-	-	-	-	-
7	2	-	2	2	-	-	-	-	1	1	1	1	1	1
8	1	-	1	1	-	-	-	1	-	1	1	1	-	-
9	1	-	1	1	-	-	-	1	1	-	1	1	-	-
10	1	-	1	1	-	-	-	1	1	1	-	1	-	-
11	1	-	1	1	-	-	-	1	1	1	1	-	-	-
12	1	-	1	1	-	-	-	1	-	-	-	-	-	1
13	1	-	1	1	-	-	-	1	-	-	-	-	1	-

IV. Initial Exploration

To identify relationships between ingredients a multi-level clustering algorithm was applied to the ingredient co-occurrence graph. The coarsening algorithm used is described in [4] which utilizes the modularity matrix

to hierarchically aggregate vertices by maximizing the modularity. Modularity of a clustered graph is a scalar that represents the quality of the clustering. High modularity means there are many edges between vertices within the aggregates and few edges between the aggregates themselves. When any two vertices are joined to make a new aggregate, all the edges are brought along and if both had edges to a mutual vertex the new weight to the mutual vertex is the sum of those edges.

Following is an overview of the algorithm to select vertices to merge to maximize the modularity as described in [4].

The ingredient co-occurrence matrix is a symmetric $n \times n$ matrix $A = (a_{ij})$. First, the rowsums of A are calculated. These rowsums make $r \in \mathbb{R}^n$.

$$r_i = \sum_j a_{ij} \tag{1}$$

Let the total rowsum be

$$T = \sum_i r_i \tag{2}$$

And the normalized rosums

$$\alpha_i = \frac{r_i}{T} \tag{3}$$

The modularity matrix, $B = (b_{ij})$, is defined

$$B = A - \frac{1}{T}rr^T \tag{4}$$

Now consider a non-overlapping partition of the vertex set $\{\mathcal{A}\}$. The new normalized rowsum and aggregate edges of each aggregate, $\mathcal{A}$

$$\alpha_{\mathcal{A}} = \frac{1}{T}\sum_{i \in \mathcal{A}} r_i \tag{5}$$

$$a_{\mathcal{A}\mathcal{B}} = \sum_{i \in \mathcal{A}}\sum_{j \in \mathcal{B}} a_{ij} \tag{6}$$

The modularity of the partition, $\mathcal{Q}$, is

$$\mathcal{Q} = \frac{1}{T}\sum_{\mathcal{A}}\sum_{i \in \mathcal{A}}\sum_{j \in \mathcal{A}} b_{ij} \tag{7}$$

The goal of the algorithm is to select the next partitioning to maximize the change of modularity. To do so we create the change of modularity matrix, $\Delta\mathcal{Q}$. Both B and $\Delta\mathcal{Q}$ have the desirable property of zero rowsums. Consider merging aggregates $\mathcal{A}$ and $\mathcal{B}$ from $\{A\}$. The resulting change in modularity, $\Delta\mathcal{Q}_{\mathcal{A}\mathcal{B}}$

$$\Delta\mathcal{Q}_{\mathcal{A}\mathcal{B}} = 2\left(\frac{a_{\mathcal{A}\mathcal{B}}}{T} - \alpha_{\mathcal{A}}\alpha_{\mathcal{B}}\right) \tag{8}$$

Using $\Delta\mathcal{Q}$ we can find which aggregate pairs should be merged by looking for the maximum value of each row. Construct the vector, $p = (p_i)$, which contains the maximum positive value from $\Delta\mathcal{Q}_{ij}$ for that row. If no positive values exist then $p_i = -1$.

$$p_i = \underset{j:i \neq j}{argmax}\{\Delta\mathcal{Q}_{ij}|\Delta\mathcal{Q}_{ij} > 0\} \tag{9}$$

This p vector can be used to construct the piecewise constant interpolation matrix, $P = (p_{ij})$. All pairs where $p_i = j$ and $p_j = i$ will be merged to create A_c. P is a sparse $n \times n_c$ matrix where each column is a aggregate in the coarsened graph. These columns will have all 0's with 1 at each i in the aggregate. This gives us

$$A_c = P^T A P \tag{10}$$

$$\Delta\mathcal{Q}_c = \frac{2}{T}A_c - \frac{2}{T^2}r_c r_c^T \tag{11}$$

The efficiency of this algorithm comes from exploiting the sparse triple matrix product $P^T A P$ and matrix vector products to construct $r_c = P^T r$, which can both be calculated in parallel.

Steps 9 through 11 are repeated to give a hierarchy of partitions. The algorithm completes when there is no positive values in $\Delta\mathcal{Q}$. Note that because $\Delta\mathcal{Q}$ has 0 rowsums, it is a weighted graph Laplacian and the modularity, $\mathcal{Q}$, is the sum of the diagonal. The algorithm can be terminated early if the coarsening factor isn't below a desired value. If A is $n \times n$ and A_c is $m \times m$ the coarsening factor is $\frac{m}{n}$. Other heuristics can be applied to speed the coarsening such as increasing the number of passes when constructing the p joining vector, starting with leaf aggregation at each level, and applying random weights to $\Delta\mathcal{Q}$ at each level.

To visualize the structure of this network these partitions were embedded into a three dimensional space using a multilevel approach. add more details about embedding process.

When processing an entire database with this coarsening algorithm, little structure was identified and most of the vertices aggregated to these hubs. This is expected as certain ingredients such as sugar, butter, and flour are extremely common. When a database was filtered by tags and cuisine clear structure was identified. Some clusterings such as filtering by pasta, salad, and cocktail result in tightly grouped clusters that contain intuitive groups of ingredients demonstrating there is clear structure in the network.

V. Recipe Generation

This section mostly describes where I am currently at and problems I am encountering.

Autoencoders [5] and Generative Adversarial Nets (GAN) [6] are well known for their capabilities to algorithmically generate data. These are both neural network infrastructures used to produce novel features in some vector space.

Fig. 1. Embedding of ingredients in pasta recipes

Fig. 2. Embedding of ingredients in salad recipes

GAN work by using two components - a discriminator and a generator. The generator takes in some random noise and outputs some data. The discriminator is a binary classifier that determines if the input is real or generated. Both networks start untrained and the discriminator is shown some examples of real data and gradient descent is applied to update the weights. Then it is given a batch of the generated data and its weights are updated again. The weights for the generator are updated based on the output of discriminator. This provides a game mechanism where the two networks attempt to outperform one another.

An autoencoder is an encoder-decoder infrastructure that attempts to perform dimensionality reduction in the encoder and reconstruct the input in the decoder. It's loss is calculated during training based on how far off the output is from the input. This is an appealing design for my application because the encoder and decoder

components can be interchanged. As an example, I could train an autoencoder on my dataset filtered by southeast Asian cuisine and train another on southwestern US cuisine and attempt to put a southeast Asain spin on a radically different cuisine. The other advantage with this infrastructure is that it is quite simple and easy to implement and tune hyperparameters.

The primary issue for using neural networks for data generation is developing a quality encoding for the training set. Using the ingredient co-occurrence column vectors as encodings for ingredients is a one hot encoding of the data. One hot encodings do not perform well in sparse high cardinality settings such as this. To attempt dimensionality reductions SVD was applied to the ingredient co-occurrence matrix and the first 10 eigenvectors were preserved from the U matrix to create dense low cardinality features for each ingredient. A popular implementation of GAN utilizes convolutional neural networks and are known as DCGAN [7]. These have been shown to work well for generation of realistic images of faces and objects.

References

[1] "Food services and drinking places: Naics 722," https://www.bls.gov/iag/tgs/iag722.htm, accessed: 2020-04-06.

[2] B. P. Majumder, S. Li, J. Ni, and J. McAuley, "Generating personalized recipes from historical user preferences," in Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP). Hong Kong, China: Association for Computational Linguistics, nov 2019, pp. 5976–5982. [Online]. Available: https://www.aclweb.org/anthology/D19-1613

[3] M. Lynch, "Ingredient phrase tagger," https://github.com/aujxn/ingredient-phrase-tagger, accessed: 2020-02-22.

[4] B. G. Quiring and P. S. Vassilevski, "Properties of the graph modularity matrix and its applications," Lawrence Livermore National Lab, Livermore, California, Tech. Rep. LLNL-TR-779424, 6 2019.

[5] A. Makhzani, J. Shlens, N. Jaitly, I. Goodfellow, and B. Frey, "Adversarial Autoencoders," arXiv:1511.05644 [cs], May 2016, arXiv: 1511.05644. [Online]. Available: http://arxiv.org/abs/1511.05644

[6] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative Adversarial Nets," in Proceedings of the 2014 Conference on Advances in Neeural Information Processing Systems (NIPS), June 2014, pp. 2672–2680.

[7] A. Radford, L. Metz, and S. Chintala, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks," in Proceedings of the 2016 International Conference on Learning Representations (ICLR), Jan. 2016, arXiv: 1511.06434. [Online]. Available: http://arxiv.org/abs/1511.06434